
Channel-Aware Training, Not Closed-Loop Control, Drives Robust Regeneration in Neural Cellular Automata

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 We set out to test whether closed-loop control of global chemical signals makes
2 regenerative Neural Cellular Automata (NCAs) more robust, and built the evalu-
3 ation to prove it: recurring multi-block lesions larger than the perception radius,
4 a small evolved controller for three broadcast modulator channels, and held-out
5 damage seeds. On a single trained model, the experiment looked like a success.
6 However, crossing five independently trained parent seeds (two independent con-
7 troller evolutions each; preregistered) dissolves that success into three findings.
8 First, the robustness we had attributed to control is mostly a property of training:
9 parents trained with channels present beat unmodulated siblings in five of five seeds
10 (median final-Hamming reduction 0.008, up to ≈ 0.14 on the most fragile parent)
11 with modulation pinned to neutral. Second, evolved controllers add little on top: on
12 four parents the effect is absent or noise-level; on the fragile parent both evolutions
13 reproducibly help (0.035–0.038 $\rightarrow$ 0.029–0.030, survival 0.975 $\rightarrow$ 1.00); yet
14 its tonic vector is nearly identical (cosine 0.99) to a sibling’s that does not help,
15 so the benefit is an interaction with parent dynamics, not a distinctive operating
16 point. All ten controllers emit flat, nonzero, parent-specific constants with no
17 lesion-locked response. Third, these constants are organism-locked: a single-donor
18 probe penalizes five of five siblings (lethally in two), a full 5×5 transfer matrix
19 shows most transfers harmful and none better than the recipient’s own controller,
20 and injecting each donor’s tonic constant directly, no controller at all, reproduces
21 its transfer outcomes in 46 of 50 cells, including all lethal ones. Single-parent
22 evaluation would have approved all of it, making the attribution protocol itself the
23 transferable artifact.

24 1 Introduction

25 Growing Neural Cellular Automata (GNCAs) grow a target morphology from a single seed and
26 regenerate it after damage using one local, differentiable update rule [1]. The catch is perception:
27 a cell sees only its immediate neighborhood, so wounds larger than a few cells contain tissue with
28 no signal to integrate, and a severed fragment has no way to know what it was part of. The natural
29 remedy is a global signal. Prior work, our own single-parent study included (Appendix H), adds
30 global modulator channels and evolves release policies for them, reporting the outcome on one trained
31 model.

32 That single-model habit is the problem this paper is about. When a modulated NCA beats an
33 unmodulated one, the channels being *present during training*, the controller *acting at run time*, and
34 whatever *transfers* to another organism are confounded. Our single-parent study learned this the hard
35 way: it credited closed-loop control for a $2 \times$ improvement that vanished under cross-parent testing,

where the controller transferred lethally to siblings while zero-modulation channel parents beat the unmodulated baseline in every seed (Appendix C). The improvement was real, but the attribution was wrong.

We therefore ran the preregistered attribution study that the original design needed with primary statistic and decision thresholds fixed before data collection. Five parent seeds, each trained from scratch as a $K=0$ pair and a channel-aware $K=3$ pair; *two* independent controller evolutions per channel parent; four conditions per seed, all on held-out damage seeds; then, post-hoc, a full 5×5 transfer matrix and a tonic-transplant condition (Section 3; preregistrations and analysis scripts timestamped before the runs).

Contributions.

1. A preregistered robustness-attribution protocol for regenerative NCAs: adversarial damage calibrated to exceed perception, crossed with parent-seed variation, decomposed into E_{train} , E_{ctrl} , E_{transfer} , with calibration probes that keep the evolution landscape interpretable (Appendix F).
2. The attribution, settled causally: training with channels is the robust cause (5/5 seeds); the controller effect is parent- dependent (absent or noise-level on four parents, reproducible on the fragile one); the evolved artifact is a parent-specific tonic constant that penalizes every sibling it meets (lethally in two), and injecting the constant alone reproduces its transfer outcomes: the matrix (23/40 harmful, none beating the recipient’s own) closes the case.

2 Related Work

Neural cellular automata. GNCAAs grow patterns from a seed and regenerate after damage [1], with extensions to self-classification [2] and textures [3]. All rely on local perception, and regeneration degrades beyond the perception radius; the regime that our benchmark exploits. Signal channels [4], goal conditioning [5], and information-dynamical analyses of self-maintenance [6] all treat the signal as fixed input or emergent byproduct instead of a controlled output; we ask what is actually helping when such a channel *appears* to help.

Search over evaluation regimes. Our design follows a concern shared by novelty search [7] and quality-diversity optimization [8]: a fixed objective can misdirect search and hide failure modes. Environment-generation methods co-evolve challenges with solutions. POET [9], PAIRED [10], ACCEL [11]; Our benchmark does not yet evolve the damage distribution and serves as the stationary baseline such co-evolution can be judged against.

3 Experimental design

Damage regime. Recurring multi-block lesions; every 150 steps, $n=4$ contiguous 16×16 blocks at seeded positions, $T=2000$ (the information-isolation regime of Appendix E). Damage seeds 0–7 drive evolution; held-out seeds 10000–10007 drive all reported numbers; the schedule is deliberately adversarial (Appendix F).

Parents and controllers. Per seed s we train a $K=0$ and a channel-aware $K=3$ parent from scratch (Appendix D); the $K=3$ parent’s three global modulator channels are the only non-local pathway, each carrying a *tonic* (slow baseline) and *phasic* (event-triggered, fast-decaying) component: the distinction from biological neuromodulation. A 259-parameter controller ($4 \rightarrow 32 \rightarrow 3$, tanh) reads four target-free grid statistics and sets release every 10 steps; each parent gets *two* independently evolved controllers via CMA-ES [12] in Evosax [13] (population 64, $\sigma_0=0.01$, 300 generations, different evolution seeds, event-weighted Hamming objective; Appendix G).

Preregistered comparison. All conditions run on the same held-out damage seeds (5 condition seeds $\times$ 8 damage seeds) from a shared $t=0$ state grown by the $K=3$ parent. The primary statistic is $\Delta_s = H(K=3, m=0, s) - H(K=3, \text{own ctrl}, s)$, positive when the evolved controller helps *its* own parent. Fixed before data collection: *controller effect supported* if $\Delta_s > 0$ in $\geq 4/5$ seeds with sign-consistent differences; *channel-training effect supported* if $H(K=0, s) - H(m=0, s) > 0$ in $\geq 4/5$;

transfer failure supported if the July controller underperforms the own-controller substantially in $\geq 4/5$. Effects are reported per seed (median and range); no significance claims at five seeds. The defense study reported here fixed the two-evolutions-per-parent rule in its preregistration before launch (experiment_results/20260818_voseed_defense/PREREGISTRATION.md in the repository history). Controller-output(m_t) series distinguish tonic calibration from event-locked policy (Appendix B). A post-hoc follow-up evaluates the full 5×5 transfer matrix among the defense parents (both controller replicas, 8 damage $\times$ 3 condition seeds per cell; Table A3) plus a *tonic transplant* injecting each donor’s realized mean m_t as a constant.

4 Results

Table 1 reports the preregistered decomposition; the full per-seed, per-condition numbers behind it are in Appendix A.

Table 1: Final Hamming (mean over 5 condition seeds $\times$ 8 held-out damage seeds). Survival in parentheses where it departs from 1.00. Training helps in 5/5 seeds; the controller effect is parent-dependent.

Parent seed	$K=0$	$K=3, m=0$	Own ctrl (2 runs)	Interpretation
0	0.035–0.036	0.029	0.029–0.030	channel effect; no control effect
1	0.163–0.175 (surv. 0.05–0.08)	0.035–0.038 (surv. 0.975)	0.029–0.030 (surv. 1.00)	channel effect + replicated tonic benefit
2	0.030	0.018–0.021	0.022–0.026	channel effect; controller slightly worse
3	0.049–0.052	0.028–0.029	0.032–0.033	channel effect; controller slightly worse
4	0.042–0.045	0.033	0.034–0.035	channel effect; no control effect

Outcome. *Channel-training supported* (positive in 5/5 seeds; bar $\geq 4/5$); *transfer failure supported* (penalty 5/5, lethal in 2). Under the preregistered rule the controller effect is classified as *parent-dependent*, preregistered Outcome C: absent or noise-level on four parents, reproducibly beneficial on the fragile one (both evolutions: 0.035–0.038 $\rightarrow$ 0.029–0.030, survival 0.975 $\rightarrow$ 1.00).

What the controller actually emits. All ten evolved controllers emit a *constant at a nonzero level*: within-rollout per-channel std of m_t is 0.0002–0.0046, tonic with no lesion-locked response in 10/10 runs, while channel means sit at parent-distinct offsets (Figure A1). Evolution neither collapses to neutral output nor discovers a policy: it finds a parent-specific tonic calibration (whose predictive limits are discussed in Section 5).

Cross-parent transfer matrix. The full 5×5 matrix among the defense parents, both replicas, three condition seeds per cell, 40 off-diagonal cells (Table A3, Figure 1), 23/40 transfers are harmful (survival < 0.9 or final Hamming worse than the recipient’s zero-output baseline by > 0.005), 15 are indistinguishable from zero-output, and 2 beat it. No foreign controller beats the recipient’s own controller beyond noise. Lethal transfer replicates across controller replicas and condition seeds: donor s3 kills recipient s2 in both replicas (survival 0.00); donor s0 on recipient s1 is lethal in e1 and near-lethal in e2 (Figure 1A).

Tonic alignment structures transfer. All lethal instances pair strongly negative tonic-vector cosines (s3 $\rightarrow$ s2: -0.61 ; s0 $\rightarrow$ s1: -0.93), while the only beneficial transfers, beating zero-output and matching the own controller, occur exclusively within the tonic-aligned pair s4 $\rightarrow$ s1 (cosine $+0.99$, both replicas; adjacent to the diagonal in Figure 1). Across all 40 cells, tonic cosine correlates with transfer penalty ($r = -0.30$) and survival




Figure 1: Transfer redesigned for attribution (replica-averaged; per-cell Δ from the recipient’s own controller; survival as the split lower triangle; rows/columns cosine-ordered so the aligned pair s4–s1 borders the diagonal and the anti-aligned pair s3–s2 sits at opposite ends). Panel A: controllers. Panel B: the same matrix with each donor’s tonic constant injected instead, matching A in 46/50 cells, every lethal one included. Values: Table A3.

($r = +0.34$); alignment is not sufficient, as some benign cells carry cosines to -0.95 .

Tonic transplant. Injecting each donor’s realized mean m_t as a constant in place of its controller (both replicas, three condition seeds) reproduces the full controller’s outcome in 36/40 off-diagonal cells, including every lethal transfer (the four mismatches are magnitude differences inside already-lethal cells), and matches the recipient’s own controller in all 10 diagonal cases (46/50 overall; Appendix A). The transferable component is thus the tonic setpoint; the controller’s small dynamic residual did not measurably contribute, on-parent or off.

5 Discussion

Which component causes robustness. Channel-aware training, not run-time control. Parents that grew up with global channels beat their unmodulated siblings on every seed; the channels are a developmental scaffold, making the organism more robust even when they carry no signal at run time. What evolution adds is narrow: nothing measurable on four parents; on the fragile parent, a final stabilization that both independent evolutions find, and since that parent’s tonic vector is nearly identical to a sibling’s that gains nothing from it, the benefit is not a property of the tonic. We conjecture the fragile parent’s dynamics sit near a stability boundary, so a small tonic push re-stabilizes the attractor (unverified). The broader lesson: held-out *damage* seeds approved every condition above; crossing *parent* seeds is what separated training from control, benefit from harm, and policy from constant.

Why cross-parent transfer fails, and when it does not. The full matrix (Table A3) sharpens the single-donor probe: no foreign controller beats the recipient’s own controller, most transfers are harmful (23/40), and the only beneficial ones stay within the tonic-aligned pair (s4→s1). The tonic transplant makes it causal: what transfers, for good or ill, is the constant (36/40 off-diagonal, 10/10 diagonal, every lethal cell included). Each controller’s output is a calibration against its own parent’s channel weights, so the same release level drives different dynamics in a sibling; evolution found an operating point for one organism, not a modulation law.

Limitations and future work. Five parents, two evolutions each; one morphology; one damage family; the matrix covers only these parents and replicas; with one controller-responsive parent in five, their prevalence is unknown. Next: multi-parent evolution and non-stationary, diversity-driven damage co-evolution, for which this benchmark is the controlled baseline.

Conclusion. In regenerative NCAs under adversarial recurring damage, robustness comes from channel-aware training (5/5 seeds); evolved modulation is a parent-dependent tonic calibration, noise-level on most parents, beneficial on the fragile one, parent-locked under transfer. Attribute robustness across model seeds, not only damage seeds.

References

- [1] A. Mordvintsev, E. Randazzo, E. Niklasson, and M. Levin. *Growing Neural Cellular Automata*. Distill 5(2):e23, 2020.
- [2] E. Randazzo, A. Mordvintsev, E. Niklasson, M. Levin, and S. Greydanus. *Self-classifying MNIST Digits*. Distill, 2020.
- [3] E. Niklasson, A. Mordvintsev, E. Randazzo, and M. Levin. *Self-Organising Textures*. Distill 6(2), 2021.
- [4] J. Stovold. *Neural Cellular Automata Can Respond to Signals*. ALIFE 2023; arXiv:2305.12971.
- [5] S. Sudhakaran, E. Najarro, and S. Risi. *Goal-Guided Neural Cellular Automata: Learning to Control Self-Organising Systems*. arXiv:2205.06806, 2022.
- [6] A. Masumori, H. Sato, and T. Ikegami. *Structured Fluctuations and the Information Dynamics of Self-Maintenance in Growing Neural Cellular Automata*. arXiv:2607.12403, 2026.
- [7] J. Lehman and K. O. Stanley. *Abandoning Objectives: Evolution Through the Search for Novelty Alone*. Evolutionary Computation, 19(2):189-223, 2011.
- [8] J.-B. Mouret and J. Clune. *Illuminating Search Spaces by Mapping Elites*. arXiv:1504.04909, 2015.
- [9] R. Wang, J. Lehman, J. Clune, and K. O. Stanley. *Paired Open-Ended Trailblazer (POET): Endlessly Generating Increasingly Complex and Diverse Learning Environments and Their Solutions*. arXiv:1901.01753, 2019.
- [10] M. Dennis, N. Jaques, E. Vinitzky, A. Bayen, S. Russell, A. Critch, and S. Levine. *Emergent Complexity and Zero-Shot Transfer via Unsupervised Environment Design*. NeurIPS 2020.
- [11] J. Parker-Holder, M. Jiang, M. Dennis, M. Samvelyan, J. Foerster, E. Grefenstette, and T. Rocktäschel. *Evolving Curricula with Regret-Based Environment Design*. ICML 2022; arXiv:2203.01302.
- [12] N. Hansen. *The CMA Evolution Strategy: A Comparing Review*. In *Towards a New Evolutionary Computation*, Studies in Fuzziness and Soft Computing vol. 192, pages 75-102. Springer, 2006.
- [13] R. T. Lange. *evosax: JAX-Based Evolution Strategies*. arXiv:2212.04180, 2022.

A Full per-seed, per-condition results

Tonic transplant. Replacing each donor controller by its realized mean m_t injected as a constant (same protocol, both replicas, 3 condition seeds) reproduces the controller’s outcome in 36/40 off-diagonal cells and

Table A1: Final Hamming (mean $\pm$ SD, 5 condition seeds $\times$ 8 held-out damage seeds) and survival for each of the ten runs (5 parent seeds $\times$ 2 independent controller evolutions) under hard recurring multi-block damage, $T=2000$. Conditions: $K=3$ parent with modulation pinned to zero; the parent’s own evolved controller; the $K=0$ parent.

Run	$K=3, m=0$		$K=3, \text{own ctrl}$		$K=0$	
	H	surv	H	surv	H	surv
s0 _e 1	0.0290 $\pm$.0135	1.000	0.0303 $\pm$.0134	1.000	0.0361 $\pm$.0137	1.000
s0 _e 2	0.0294 $\pm$.0144	1.000	0.0292 $\pm$.0136	1.000	0.0354 $\pm$.0109	1.000
s1 _e 1	0.0384 $\pm$.0266	0.975	0.0286 $\pm$.0177	1.000	0.1628 $\pm$.0389	0.075
s1 _e 2	0.0353 $\pm$.0256	0.975	0.0301 $\pm$.0210	1.000	0.1746 $\pm$.0356	0.050
s2 _e 1	0.0183 $\pm$.0145	1.000	0.0225 $\pm$.0149	1.000	0.0299 $\pm$.0178	1.000
s2 _e 2	0.0208 $\pm$.0170	1.000	0.0258 $\pm$.0161	1.000	0.0296 $\pm$.0174	1.000
s3 _e 1	0.0281 $\pm$.0177	1.000	0.0332 $\pm$.0165	1.000	0.0491 $\pm$.0084	1.000
s3 _e 2	0.0289 $\pm$.0182	1.000	0.0321 $\pm$.0164	1.000	0.0525 $\pm$.0136	1.000
s4 _e 1	0.0332 $\pm$.0233	1.000	0.0336 $\pm$.0223	1.000	0.0420 $\pm$.0111	1.000
s4 _e 2	0.0332 $\pm$.0225	1.000	0.0345 $\pm$.0222	1.000	0.0452 $\pm$.0093	1.000

Table A2: Cross-parent transfer probe (unchanged from the per-parent study): the July donor controller evaluated on each recipient parent. Final Hamming and survival on held-out damage seeds; transfer is a penalty in 5/5 recipients, lethal in 2.

Recipient seed	Final Hamming	Survival
0	0.036	1.00
1	0.249	0.00 (lethal)
2	0.434	0.00 (lethal)
3	0.100	0.45
4	0.041	1.00

205 10/10 diagonal cells (46/50 overall; Figure 1B), including every lethal
206 cell, where the transplant kills exactly as the controller does. The
207 four mismatches are magnitude differences inside already-lethal cells.
208 Under the tested regime, the controllers’ small dynamic residual did not
209 measurably contribute, on-parent or off (full CSVs in the repository).
210 AUC mirrors the same ordering in every seed (full CSVs in the repository).

211 B Controller-output (m_t) diagnostics




Figure A1: Controller output m_t over a full evaluation rollout, per parent seed, both independent evolutions overlaid (solid/dashed), one line per modulator channel. After a brief initial transient while the perceptual state settles, every controller holds a constant, parent-distinct offset through all lesion events: tonic calibration, not an event-locked policy. Note that the two independent evolutions of the same parent sometimes settle at slightly different setpoints.

212 The tonic vector alone does not predict benefit: seeds 1 and 4 emit nearly
213 identical tonic vectors (cosine similarity 0.993, Euclidean distance 0.0065),
214 yet only seed 1 benefits from its controller. The benefit arises from the

Table A3: Full 5×5 cross-parent transfer matrix among the five defense parents, both controller replicas (top: e1; bottom: e2). Rows: donor controller; columns: recipient parent. Cells give mean final Hamming / worst-case survival over 3 condition seeds $\times$ 8 held-out damage seeds (24 rollouts per cell). Diagonal (bold): recipient’s own controller (survival 1.00 throughout). † lethal (survival 0.00). Across the 40 off-diagonal cells: 23 harmful, 15 indistinguishable from zero-output, 2 beneficial; the beneficial cells are exclusively $s4 \rightarrow s1$, the tonic-aligned pair (cosine +0.99), in both replicas.

e1 donor	recip. s0	recip. s1	recip. s2	recip. s3	recip. s4
s0	0.0315	0.470/0.00 †	0.097/0.38	0.028/1.00	0.052/0.88
s1	0.035/1.00	0.0306	0.023/1.00	0.049/1.00	0.035/1.00
s2	0.029/1.00	0.043/0.88	0.0201	0.033/1.00	0.092/0.50
s3	0.035/1.00	0.036/0.88	0.208/0.00 †	0.0295	0.044/1.00
s4	0.034/1.00	0.030/1.00	0.023/1.00	0.044/1.00	0.0340
e2 donor	recip. s0	recip. s1	recip. s2	recip. s3	recip. s4
s0	0.0291	0.108/0.25	0.056/0.88	0.032/1.00	0.039/1.00
s1	0.032/1.00	0.0292	0.025/1.00	0.049/1.00	0.032/1.00
s2	0.041/1.00	0.074/0.62	0.0219	0.033/1.00	0.113/0.25
s3	0.028/1.00	0.058/0.75	0.239/0.00 †	0.0262	0.028/1.00
s4	0.030/1.00	0.030/1.00	0.020/1.00	0.033/1.00	0.0336

Table A4: Tonic output of the evolved controllers: per-channel mean of m_t , averaged over the two independent evolutions per seed. All ten controllers sit at nonzero, parent-distinct offsets; within-rollout per-channel std of m_t is 0.0002–0.0046 everywhere: flat, tonic output with no lesion-locked response in 10/10 runs.

Seed	mean(m) per channel (avg. of 2 runs)		
0	−0.021	+0.002	+0.027
1	+0.030	+0.012	−0.026
2	−0.026	−0.032	−0.015
3	+0.001	+0.032	−0.010
4	+0.028	+0.007	−0.025

interaction between a tonic calibration and parent-specific local dynamics,
not from a distinctive chemical operating point.

C Pilot: single-parent evaluation hides parent-locking

A three-parent pilot (2026-08-16) first exposed the confound. The single-parent controller from our original five-condition study transferred lethally to two of three sibling parents (survival 0.00, final Hamming 0.249 and 0.434), while zero-output channel parents beat the $K=0$ baseline in 3/3 seeds (0.029/0.034/0.020 vs 0.034/0.177/0.028). The original table’s headline gap was therefore part parent-training effect and part parent-seed luck, motivating this study.

D Model and training details

The grid state $x_t \in [0,1]^{96 \times 96 \times 16}$ holds 16 channels; the last four are RGBA with alpha last. Perception uses three fixed kernels (identity, Sobel- x , Sobel- y) per channel, giving $p_t \in \mathbb{R}^{48}$:

$$p * t = (K * \text{id}_x * t, K * \text{S}_{-x} * t, K * \text{S}_{-y} * t). \quad (1)$$

The update MLP maps $(48+K)$ inputs to 128 hidden units (ReLU) and back to 16 channel increments, with the final layer zero-initialized. Cells update with probability 0.5; a cell is alive when its alpha exceeds 0.1 within its 3×3 max-pooled neighborhood; dead cells are masked to zero.

233 Tonic and phasic modulator dynamics:

$$m * t^{(\text{tonic})} = \alpha m * t - 1^{(\text{tonic})} + (1 - \alpha) c * t, \quad m_t^{(\text{phasic})} = m * t - 1^{(\text{phasic})} \cdot e^{-\Delta t / \tau}, \quad (2)$$

234 with $\alpha = 0.95$, $\tau = 20$, and the injected level the clipped sum, broadcast to
 235 all cells and concatenated to perception (input dim $48 + 3 = 51$).

236 Parents train on a 40×40 emoji target zero-padded to the canvas, with
 237 pool-based training (pool 1024, batch 8, worst-in-batch reseeding), random
 238 square-erasure damage mid-episode, MSE loss on RGBA, and Adam at 10^{-3} for
 239 8000 steps. A first attempt at 2×10^{-3} diverged to NaN near step 5532 via a
 240 late-stage loss-spike cascade; 10^{-3} absorbs the identical spike.

241 E Perception-radius sweep

242 A single-lesion sweep (radii $\{2, 4, 8, 16\}$, single- and multi-site, with and
 243 without modulator channels) maps the boundary of local repair. Recovery
 244 is near-perfect for small lesions and degrades once lesions exceed the
 245 effective perception radius. Two failure signatures motivate this work:
 246 severed fragments that survive but cannot reattach, die, or decide what to
 247 grow into, and persistent half-alpha scar tissue at wound sites. Both are
 248 information failures: the misbehaving cells are exactly the cells cut off
 249 from global context.




Figure A2: E1 lesion sweep. Final Hamming distance as a function of lesion radius and kind, with and without modulator channels (mean $\pm$ SD over 5 damage seeds). Inset: post-regrowth debris and scar tissue.

250 F Calibration probes

251 **Collapse probe.** The target mask is only 4.6% alive, so a fully-dead
 252 grid scores Hamming ≈ 0.046 , close to a struggling organism's ≈ 0.02 .
 253 We verified empirically that the all-dead state scores 0.0461 against
 254 struggling-neutral's 0.0205: death loses by $2.2\times$, so the alive-fraction
 255 floor is disabled (0.0). Notably, the runbook-default floor of 0.3 would
 256 have *penalized every healthy candidate* (a correct lizard lives at ≈ 0.057
 257 alive), pushing evolution toward overgrowth, a case where calibrating from
 258 the experiment's own data overrode a plausible default.

259 **Initial step-size probe.** Perturbing the neutral controller by Gaussian noise
 260 at increasing σ (25 samples each, hard-regime evaluation) showed: $\sigma \leq 0.01$
 261 stays healthy (some samples beat neutral); $\sigma = 0.05$ is bimodal; $\sigma = 0.3$
 262 lands *every* sample in saturated overgrowth (alive $\rightarrow$ 0.96). A first
 263 evolution run with the library-default $\sigma_0 = 0.3$ froze at fitness 0.150
 264 from generation 2 onward, $7\times$ worse than doing nothing, because CMA-ES was
 265 ranking overgrowth-degree noise. $\sigma_0 = 0.01$ brackets the neutral point and
 266 the same budget converged smoothly (Figure A3).




Figure A3: CMA-ES fitness over 300 generations ($\sigma_0=0.01$). Dashed line: neutral controller (0.0205). Best evolved fitness: 0.0135, a 34% improvement with no stall.

267 G Evolution objective

268 We evolve the controller with CMA-ES [12] (Evojax implementation [13]),
 269 population 64, initial step size $\sigma_0=0.01$, 300 generations, neutral
 270 (all-zero) controller as the initial mean. Fitness is the event-weighted
 271 mean Hamming distance to the target alpha mask over full rollouts on the
 272 eight train damage seeds:

$$f(\theta) = \frac{1}{\sum * t w_t} \sum * t = 1^T w * t H_t(\theta), \quad w_t = \exp\left(-\frac{t - \tau * \text{last}(t)}{\tau * w}\right), \quad (3)$$

273 where $\tau * \text{last}(t)$ is the most recent lesion step and $\tau_w = 150/3 = 50$. The
 274 recency kernel resets at each lesion, so slow repair is expensive even
 275 when endpoint Hamming is identical, so it de-saturates the metric and
 276 prices repair speed directly. Fitness is not the reported metric: all
 277 results report unweighted trajectory statistics on held-out seeds. Two
 278 calibration controls make this landscape interpretable: a collapse probe
 279 (verifying the objective cannot be gamed by alpha suppression) and an
 280 initial step-size ablation (showing $\sigma_0=0.3$ saturates the first population
 281 in overgrowth where no repair signal exists), both detailed in Appendix F.

282 H Single-parent five-condition study

283 Five conditions share the same protocol, horizon, and damage seeds:
 284 closed-loop (evolved controller); static (the evolved controller evaluated
 285 once at $t=0$ and held constant, isolating temporal modulation from the
 286 evolved tonic level itself); constant (fixed tonic, grid-searched over
 287 $\{-1, -0.5, 0, 0.5, 1\}$ on train seeds); random (uniform $[-1, 1]$ each step,

288 actuation-matched noise); and no modulation ($K=0$ parent). All experiments
289 ran on a single rented A100; parent training took ≈ 20 min/model, evolution
290 ≈ 2.5 h, total project cost under \$15.

291 **Modulation helps substantially (RQ1).** Relative to no modulation, modulated
292 conditions recover $2.2\times$ more completely (final Hamming 0.028-0.030 vs. 0.063)
293 and sustain $\approx 4\times$ less cumulative damage (AUC 0.016-0.017 vs. 0.064). The
294 wounds exceed the perception radius, so this improvement is attributable
295 to the broadcast channel, the only non-local pathway, which repairs damage
296 that local rules cannot diagnose.

297 **A regime boundary, not a failed ablation (RQ2).** The three modulated conditions
298 are statistically indistinguishable: final Hamming spans 0.028-0.030,
299 half-life 6.4-7.2 steps, all differences within noise (pairwise effect sizes
300 < 0.15). We read this as a measurement of the *regime*, not a failure of
301 the controller: the damage process is stationary and memoryless (events
302 are i.i.d. in size, number, and position), so the optimal release policy
303 is time-invariant and a fixed tonic gain is structurally adequate. There
304 is no temporal structure for adaptive scheduling to exploit. Evolution’s
305 measured contribution is automatic discovery of the near-optimal tonic
306 level in 2.5 hours from a neutral start, without the hand search the
307 constant condition required.

308 **Random modulation is lethal.** Random actuation kills every run (survival 0.00,
309 final Hamming 0.825): pilot probes show constant levels alone swing mean
310 Hamming from 0.003 to 0.93. The modulator channels are a high-gain control
311 pathway, not a free architectural bonus; the identity of the release
312 schedule determines whether the organism lives.

313 **Metric artifacts of long-horizon evaluation (RQ3).** Two artifacts surfaced only under
314 repeated damage. First, repair half-life is defined per lesion relative
315 to each post-lesion jump; on the baseline’s chronically damaged morphology,
316 new lesions often land on already-degraded regions, producing no measurable
317 jump and scoring as zero half-life, so the baseline’s 2.7 steps does *not*
318 indicate fast repair. Second, residual Hamming in all conditions is partly
319 attributable to locomotion drift during regeneration rather than permanent
320 damage; the metric conflates spatial translation with morphological error.
321 Both are documented so that future benchmark users can avoid them.

322 I Reproducibility

323 Parents train in ≈ 20 -30 min each on one A100; the defense study ran two
324 independent 300-generation controller evolutions per parent (10 total),
325 each ≈ 2.5 -3 h on the study A100, for a total defense-study cost of $\approx \$25$
326 of rented GPU time. The transfer matrix is evaluation-only over existing
327 artifacts: 40 off-diagonal evaluations, ≈ 2 minutes of A100 time per full
328 5×5 replica matrix. Preregistration, analysis script, per-seed artifacts
329 (controllers, parents, trajectories, m_t series), and the mechanical outcome
330 output are in the repository, timestamped before the results. Code will be
331 released upon publication.